

Actionability Energy: A Tangent-Space Test for Controllable Visual Futures

Anonymous submission-ready draft

Abstract

Video and world models can imagine futures that look plausible and reach a goal, yet require local changes the robot body cannot produce. We study this failure mode through *actionability energy*: the distance between an imagined representation transition and the action-reachable tangent set induced by an embodiment action map. The central claim is that Jacobian-style action maps, such as those used for visual inverse control, can also serve as diagnostic interfaces for world-model planning. They can certify whether a future is locally executable, predict failure, and repair non-actionable trajectories by compositional inference. We implement analytic and learned local Jacobian maps on a CPU benchmark with a point mass, nonholonomic car, two-link arm, contact-switching pusher, and rendered keypoint futures. Analytic actionability is near-perfect, learned local action maps recover most of the diagnostic signal, and a learned classifier can fit the synthetic benchmark but does not explain the mechanism. Smoothness, passive likelihood, action norm, inverse-dynamics reconstruction, wrong-map residuals, and forward-inverse asymmetry proxy baselines are weaker. Robustness audits expose the hard case: discontinuous contact modes. The study is intentionally controlled and does not claim real-robot validation or large video-model finetuning; it isolates a mechanism for making generated futures honest about the body that must execute them.

1 Introduction

World models and video planners give embodied agents a tempting interface: imagine futures, score them against a goal, and execute the best one. This interface is insufficient when the future is evaluated only as a passive visual sequence. A generated transition may be smooth, likely under video data, and goal-satisfying, while still being impossible for the current embodiment to realize in one local step. A nonholonomic car cannot move sideways without changing heading. A pusher cannot move an object before contact. An arm near a singular configuration cannot realize every Cartesian displacement within its action limits.

This paper studies the local test that is missing from such futures. Let z_t be a representation state and let $J_\phi(z_t)$ map small controls into representation-space changes. The local action-reachable tangent set is

$$\mathcal{A}(z_t) = \{J_\phi(z_t)u : u \in \mathcal{U}\}.$$

For an imagined transition $\Delta z_t = z_{t+1} - z_t$, we define the actionability residual

$$r_t = \min_{u_t \in \mathcal{U}} \|\Delta z_t - J_\phi(z_t)u_t\|_W^2 + \lambda \|u_t\|^2.$$

The score is small when the transition can be explained by a bounded local action and large when the generated future leaves the embodiment’s reachable tangent set.

The idea is complementary to recent Jacobian-based robot policies and executable-video systems. Neural Jacobian Fields and VERA show that local visual Jacobians can translate desired visual change into action [7, 8]. EVA aligns video world models with inverse-dynamics rewards [12]. World Action Verifier asks whether world-model futures are executable through forward-inverse asymmetry [10]. We ask a narrower mechanism-level question: can the same local action-sensitivity object used for inverse control also certify and repair imagined futures before execution?

Contributions. First, we formalize an actionability residual for generated representation trajectories and connect it to local execution error. Second, we make learned local Jacobian maps central: the same map




Figure 1: A generated future can be visually plausible while leaving the local action-reachable tangent set. Actionability energy turns an embodiment Jacobian into a diagnostic interface: low-energy futures are executable, while high-energy futures are repaired by composed trajectory inference.

supports inverse action recovery, failure prediction, and future repair. Third, we build a CPU-first benchmark with analytic states and rendered keypoint futures, then compare against hostile baselines including inverse-dynamics reconstruction, a learned executability classifier, and a forward-inverse asymmetry proxy. Fourth, we run robustness audits over data scale, Jacobian noise, wrong maps, representation noise, and contact-mode mismatch. Fifth, we demonstrate embodiment swapping: the same imagined displacement can be feasible for one body and infeasible for another by changing only J_ϕ .

2 Related Work

Jacobian representations for robot action. Neural Jacobian Fields learn visual Jacobian fields that map desired visual changes to robot actions across diverse robots [7]. VERA uses a video planner with a Jacobian-based inverse dynamics module to produce generalist robot policies [8]. Our work does not improve NJF or reproduce VERA-scale robot experiments. It extends the role of the same local action-sensitivity object: instead of only asking which action would realize a desired visual change, we ask whether the desired change is locally reachable at all and how to repair it if not.

Executable video and world-action verification. EVA aligns video world models with executable robot actions through inverse-dynamics rewards [12]. World Action Verifier studies self-improving world models through forward-inverse asymmetry [10]. These works are the closest novelty boundary. Our residual is not a learned verifier loop or a large-video-model finetuning objective. It is a low-compute tangent-space certificate computed from analytic, finite-difference, or learned local action maps, and it is explicitly used as a repair energy.

Energy-based planning and compositional inference. Energy-based planning frames control as inference over composed objectives [5]. LEAP performs iterative energy minimization with sequence models [2], and compositional generative modeling studies how separate energy factors can be combined at inference time [4]. We use actionability as one factor in such a composed trajectory energy, rather than as a decorative name.

Queryable scene and world representations. Neural scene representations for visuomotor control demonstrate that learned representations can support robot control queries rather than only rendering [9]. RoboDreamer and Large Video Planner push world-model imagination toward robot planning [1, 13]. Our claim is narrower: a representation for embodied planning should expose local actionability, not only visual prediction or geometric state.

Failure modes in passive learning. Visual foresight and model-predictive control showed that learned video prediction can be useful for manipulation [6]. Causal confusion and observational overfitting show that passive correlations can mislead policies or value estimates [3, 11]. Actionability is a related embodied failure mode: a passive future can be plausible but physically unreachable for the current body.

3 Actionability Energy

3.1 One-Step Residual

Given a representation transition $\Delta z_t = z_{t+1} - z_t$ and local action map $J_\phi(z_t) \in \mathbb{R}^{d_z \times d_u}$, the unconstrained residual has the closed form

$$\begin{aligned} u_t^* &= (J^\top W J + \lambda I)^{-1} J^\top W \Delta z_t, \\ r_t &= \|\Delta z_t - J u_t^*\|_W^2 + \lambda \|u_t^*\|^2. \end{aligned}$$

When actions are box-bounded, $u \in [u_{\min}, u_{\max}]$, the implementation solves the same quadratic program under bounds. Because the benchmark action spaces are one- and two-dimensional, we use an exact active-set enumeration for speed and keep a generic SciPy optimizer fallback for larger action spaces.

3.2 Trajectory Energy

For a candidate future $\tau = z_{0:T}$ and goal g , we compose

$$E(\tau) = E_{\text{world}}(\tau) + \alpha E_{\text{goal}}(\tau, g) + \beta \sum_{t=0}^{T-1} r_t + \gamma E_{\text{smooth}}(\tau) + \eta E_{\text{constraints}}(\tau).$$

E_{world} keeps the repaired trajectory near the proposed future, E_{goal} pulls the terminal state toward the task target, and $\sum_t r_t$ penalizes non-actionable local transitions. In this package, repair is implemented as projected trajectory inference: build a reachable rollout from local action projections or simple embodiment-aware controllers, then blend it with world and goal terms. This is intentionally CPU-first and should be read as a benchmark solver, not as a claim of optimal nonlinear trajectory optimization.

3.3 Learning the Action Map

The analytic Jacobian is not required. We collect small exploratory rollouts (z_t, u_t, z_{t+1}) and fit a local ridge model

$$z_{t+1} - z_t \approx J_\phi(z_t) u_t.$$

The same learned map is then used for three roles: inverse action recovery, residual-based failure diagnosis, and repair. This is the Lester/NJF-facing point of the paper: a Jacobian-style embodiment representation is useful not only as an inverse-control module, but also as an interface between imagined futures and embodied feasibility.

3.4 Embodiment Swapping

The same representation displacement can be paired with different J_ϕ maps. A sideways (x, y) displacement has zero residual for a holonomic point mass but high residual for a car whose heading points forward. A pusher-object displacement is actionable only under contact. This makes actionability a property of a representation plus an embodiment, not a property of the visual future alone.

4 Theory: Plausibility Is Not Executability

Proposition 1 (Passive plausibility does not imply actionability). *There exist two trajectories τ_a, τ_b with equal smoothness and equal terminal goal reward under a passive trajectory score, such that τ_a is locally actionable and τ_b has strictly positive actionability residual.*

Proof. Consider a nonholonomic car at $z_0 = (0, 0, 0)$ with representation (x, y, θ) and local Jacobian

$$J(z_0) = \begin{bmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix}$$

after absorbing the time step into units. Let τ_a use $\Delta z = (\epsilon, 0, 0)$ and τ_b use $\Delta z = (0, \epsilon, 0)$ for $\epsilon > 0$. Both one-step trajectories can have identical Euclidean length and identical passive smoothness, and a goal-only score can assign either transition the preferred terminal reward by placing the goal at the corresponding endpoint. For τ_a , $\Delta z \in \text{Im}(J)$, so the residual can be zero up to the ridge term. For τ_b , the y component is orthogonal to $\text{Im}(J)$, so the residual is at least ϵ^2 for $W = I$ and $\lambda = 0$. Thus passive plausibility and goal reward alone do not certify actionability. $\square$

Proposition 2 (Residual lower-bounds local execution error). *Assume local dynamics in representation space satisfy*

$$z_{t+1} = z_t + J(z_t)u_t + \rho(u_t), \quad \|\rho(u_t)\| \leq c\|u_t\|^2$$

for bounded u_t . Let $r_t^0 = \min_{u \in \mathcal{U}} \|\Delta z_t - J(z_t)u\|_2^2$ be the unregularized residual. Then any one-step execution attempting to realize Δz_t has representation error at least $\sqrt{r_t^0} - c\|u_t\|^2$.

Proof. For any feasible u_t ,

$$\|\Delta z_t - (Ju_t + \rho(u_t))\|_2 \geq \|\Delta z_t - Ju_t\|_2 - \|\rho(u_t)\|_2.$$

By definition, $\|\Delta z_t - Ju_t\|_2 \geq \sqrt{r_t^0}$ for every feasible action. Applying the second-order bound on ρ gives the result. The bound is local: if actions are large or contact modes switch, the second-order term can dominate. $\square$

These propositions justify using the residual as a diagnostic rather than as a complete dynamics model. The residual detects local tangent-set mismatch; it does not prove long-horizon feasibility under arbitrary nonlinear dynamics.

5 Low-Compute Benchmark and Experiments

Environments. The benchmark uses four deterministic CPU environments. The point mass is a positive control with $J = I$. The nonholonomic car exposes sideways visual futures that are locally impossible. The two-link arm uses an end-effector Jacobian and action limits. The planar pusher has a contact-dependent Jacobian, so object motion before contact is non-actionable.

Generated futures. We use controlled future generators rather than large video models: straight-line interpolation, noisy splines, passive smoothness priors, scripted rollouts, and intentionally non-actionable futures. These generators are proxies for generated visual futures and are not claimed to match real video-model distributions.

Experiment 1: constructive counterexample. The car example in Proposition 1 is implemented directly. A sideways future has the same endpoint smoothness as a forward future under passive scoring, but the residual rejects it because the displacement lies outside the car’s local tangent set.

Experiment 2: failure prediction with hostile baselines. For each candidate, we compute residuals, recover controls, roll out the true environment, and mark failure when the executed path diverges from the imagined path. We compare analytic actionability, learned-J actionability, noisy learned-J actionability, wrong-map residual, inverse-dynamics reconstruction error, a learned classifier, a forward-inverse asymmetry proxy, smoothness, path length, action norm, passive likelihood proxy, and random score.




Figure 2: The learned action map is central, not an afterthought. Small exploratory rollouts fit J_ϕ , and the same map is used for inverse action, failure diagnosis, and future repair.




Figure 3: Representative imagined trajectories and rollouts under recovered controls. The environments isolate different reasons a plausible representation future can fail under the executing body.

Experiment 3: repair. We compare the original candidate, goal-only, world+goal, world+goal+smoothness, world+goal+actionability, full energy, learned-J full energy, noisy learned-J full energy, and wrong-map full energy. Adding actionability changes the trajectory rather than only rescoring it; it encourages futures whose local increments can be produced by bounded actions.

Experiment 4: embodiment swap. We score the same local (dx, dy) displacement under different action maps. The residual changes sharply across embodiments, showing that executability is body-conditioned.

Experiment 5: visual/keypoint futures. To avoid a purely state-space story, we render CPU keypoint representations for the car and pusher and compute residuals in keypoint space using finite-difference or learned visual Jacobians. This remains a controlled proxy rather than a large video model, but it tests the intended representation-level interface.

Table 1: Failure prediction with the actionability residual. Higher scores predict execution failure.

Environment	AUROC	AUPRC	Acc.
nonholonomic car	1.000	1.000	0.900
planar pusher	1.000	1.000	0.900
point mass	1.000	1.000	0.900
two link arm	0.990	0.966	0.811
all	0.989	0.989	0.917




Figure 4: Hostile baseline comparison. A learned classifier can pattern-fit this synthetic benchmark, while actionability provides the mechanism-level certificate. Learned-J residuals recover most of the signal; smoothness and passive likelihood can prefer non-executable futures.

6 Limitations

The benchmark is still controlled and CPU-only. It tests local representation-space feasibility, not real robot reliability. The repair solver is a projected-inference routine, not a full nonlinear optimal-control solver. The residual is local and can be misleading when second-order effects, contact discontinuities, perception errors, or long-horizon compounding dominate. Learned Jacobians are evaluated with small local regressors rather than neural Jacobian fields. The pusher exposes the central failure mode: contact discontinuities can break a smooth local map unless the robot actively probes the contact mode or uses a mode-aware representation. These limitations are intentional for a first mechanism paper, but real robot or stronger simulator validation would be required before deployment claims.

7 Conclusion

Actionability energy formalizes a simple requirement for embodied world models: imagined representation changes should lie near the local action-reachable tangent set of the body that must execute them. In controlled CPU benchmarks, analytic and learned local Jacobian maps predict execution failure, support repair, and change under embodiment swaps. The result is not a replacement for video planners, inverse dynamics, or real robot evaluation. It reframes Jacobian-style embodiment maps as diagnostic interfaces for world-model planning: they help a robot ask whether the future it imagined is actually actionable by its own body.




Figure 5: Robustness audit over learned maps. Local Jacobian maps become reliable with modest data in smooth embodiments. The contact pusher remains the hard case because the reachable tangent set changes discontinuously with contact mode.

Table 2: Repair success on diagnostic non-executable futures. The full energy includes world, goal, smoothness, and actionability terms.

Environment	Candidate	Full energy
nonholonomic car	0	1
planar pusher	0	1
point mass	0	1
two link arm	0	1

8 Reproducibility Statement

All experiments are CPU-first Python scripts. The repository includes environment code, residual solvers, repair routines, tests, raw CSV results, plotting scripts, and LaTeX source. The intended verification commands are:

```

pytest -q
python experiments/run_all.py --quick
python experiments/run_all.py --full
bash scripts/make_plots.sh
bash scripts/compile_paper.sh

```

The generated reports record command status, dependency versions, hardware, unresolved issues, and readiness level.

References

- [1] Boyuan Chen, Tianyuan Zhang, Haoran Geng, Caiyi Zhang, Peihao Li, Kiwhan Song, William T. Freeman, Jitendra Malik, Pieter Abbeel, Russ Tedrake, Vincent Sitzmann, and Yilun Du. Large video planner enables generalizable robot control, 2025. URL <https://arxiv.org/abs/2512.15840>.
- [2] Hongyi Chen, Yilun Du, Yiye Chen, Joshua B. Tenenbaum, and Patricio A. Vela. Planning with sequence models through iterative energy minimization. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=cVFD6qE8gnY>.




Figure 6: Before/after repair examples. Solid lines show imagined futures; dashed lines show rollouts from recovered controls.




Figure 7: Embodiment swap heatmaps. The same representation displacement can be actionable for one body and non-actionable for another.

- [3] Pim de Haan, Dinesh Jayaraman, and Sergey Levine. Causal confusion in imitation learning, 2019. URL <https://arxiv.org/abs/1905.11979>.
- [4] Yilun Du and Leslie Kaelbling. Compositional generative modeling: A single model is not all you need, 2024. URL <https://arxiv.org/abs/2402.01103>.
- [5] Yilun Du, Toru Lin, and Igor Mordatch. Model based planning with energy based models, 2019. URL <https://arxiv.org/abs/1909.06878>.
- [6] Chelsea Finn and Sergey Levine. Deep visual foresight for planning robot motion. In *IEEE International Conference on Robotics and Automation*, 2017.
- [7] Sizhe Lester Li, Annan Zhang, Boyuan Chen, Hanna Matusik, Chao Liu, Daniela Rus, and Vincent Sitzmann. Controlling diverse robots by inferring jacobian fields with deep networks, 2024. URL <https://arxiv.org/abs/2407.08722>.
- [8] Sizhe Lester Li, Evan Kim, Xingjian Bai, Tong Zhao, Tao Pang, Max Simchowitz, and Vincent Sitzmann. Turning video models into generalist robot policies, 2026. URL <https://arxiv.org/abs/2605.27817>.
- [9] Yunzhu Li, Shuang Li, Vincent Sitzmann, Pulkit Agrawal, and Antonio Torralba. 3d neural scene representations for visuomotor control. In *Conference on Robot Learning*, 2021. URL <https://arxiv.org/abs/2107.04004>.




Figure 8: CPU visual/keypoint futures. Actionability remains useful when z is a rendered keypoint representation rather than the raw simulator state.

- [10] Yuejiang Liu, Fan Feng, Lingjing Kong, Weifeng Lu, Jinzhou Tang, Kun Zhang, Kevin Murphy, Chelsea Finn, and Yilun Du. World action verifier: Self-improving world models via forward-inverse asymmetry, 2026. URL <https://arxiv.org/abs/2604.01985>.
- [11] Xingyou Song, Yiding Jiang, Stephen Tu, Yilun Du, and Behnam Neyshabur. Observational overfitting in reinforcement learning, 2019. URL <https://arxiv.org/abs/1912.02975>.
- [12] Ruixiang Wang, Qingming Liu, Yueci Deng, Guiliang Liu, Zhen Liu, and Kui Jia. Eva: Aligning video world models with executable robot actions via inverse dynamics rewards, 2026. URL <https://arxiv.org/abs/2603.17808>.
- [13] Siyuan Zhou, Yilun Du, Jiaben Chen, Yandong Li, Dit-Yan Yeung, and Chuang Gan. Robodreamer: Learning compositional world models for robot imagination, 2024. URL <https://arxiv.org/abs/2404.12377>.

A Additional Implementation Details

Residual solver. For unconstrained actions, the solver uses the ridge projection closed form. For bounded actions, the benchmark uses exact active-set enumeration over low, free, and high states for each action coordinate. This is exact for the convex quadratic objective in the low-dimensional action spaces used here.

Candidate generators. Candidates include true rollouts, straight-line paths, too-fast point-mass jumps, sideways car paths, Cartesian arm teleports, and object-motion-before-contact pusher paths. Labels are not assigned by construction; they are produced by recovering controls from the candidate and rolling out the environment.

Repair solver. The repair routine constructs a reachable trajectory from local action projections or embodiment-aware controllers and blends it with world, goal, and smoothness terms. The upgraded experiments evaluate analytic, learned, noisy learned, and wrong-map repair variants. This is a practical CPU solver for the benchmark. Future versions should compare against differentiable nonlinear trajectory optimization and learned video latent repair.

Visual/keypoint representation. The visual experiment renders compact keypoint states for the car and pusher. It is not a learned video generator, but it tests whether the residual can be evaluated in a representation that is closer to a visual future than raw simulator coordinates.

Novelty boundary. The closest overlap found is World Action Verifier [10]. The present work should avoid claiming to be the first executable-video verifier. Its narrower claim is a tangent-space actionability residual, benchmark, and repair energy for local embodiment-conditioned generated futures.